


Introducing zkBlob and data compression to the zkEVM

Paris, 19 July 2023 ethCC



Carlos Matallana Espinar
Protocol & Integration Team Lead

Presentation Outline



1. Intro

- Current approach
- L2 transaction cost
- EIP4844 & data compression

2. zkBlob

- Introduction
- Specification

3. Data compression

- Specification
- Examples
 - ERC20 transfer



1

Intro

Smart Contract CALL \Rightarrow

```
function sequenceBatches(BatchData[] calldata batches, address I2Coinbase)
```



batches \Rightarrow Tx_0 # Tx_1 # ... # Tx_N



RLP(nonce, gasPrice, gasLimit, to, value, data, chainId, 0, 0) # r # s # v



RLP(nonce, gasPrice, gasLimit, to, value, data, chainId, 0, 0) # r # s # v

- Smart contract to receive array of batches (sequences)
- Each batch contains transactions
- Each transaction is posted on-chain in the following format:
 - **RLP(txFields) + signature**
 - Easy to compute signed message inside the **zkevm-rom**

INTRO: L2 transaction cost

Data-availability

- Data-availability
 - paid for every byte posted on-chain

Sequencing

- Sequencing
 - constant L1 transaction cost
 - shared among all batches (up to 128 kB per smart contract call)

Prover

- Prover
 - constant hardware costs
 - for each proof computed

Aggregation

- Aggregation
 - constant L1 transaction cost
 - shared among all sequences aggregated

Summary

	\$ (%)			
	data-availability	sequencing	prover	aggregation
Ether transfer	0.238 \$ (95.78 %)	0.010 \$ (4.04 %)	0.00007 \$ (0.03 %)	0.00038 \$ (0.15 %)
ERC20 transfer	0.305 \$ (94.92 %)	0.015 \$ (4.73 %)	0.00018 \$ (0.06 %)	0.00095 \$ (0.3 %)
UniswapV2 swap	1.019 \$ (96.22 %)	0.032 \$ (3.05 %)	0.00136 \$ (0.13 %)	0.00640 \$ (0.6 %)

Cost computation

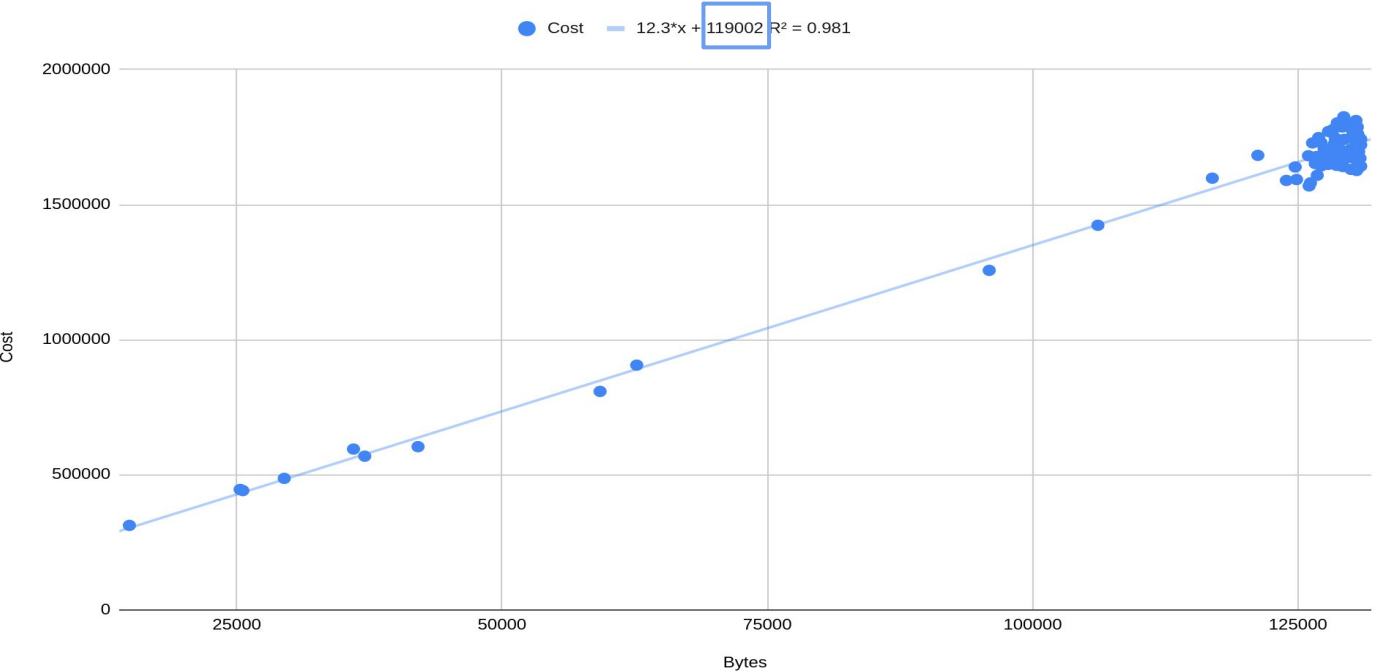
- 16 GAS byte calldata
- 128kB maximum in **sequenceBatches** SC call (geth limitation)
- Assume batches are 70% filled (**zk-counters** as larger limitant factor)
 - 350 eth transfers → 39630 bytes per batch
 - 140 ERC20 transfers → 24043 bytes per batch
 - 19 uniswapV2 swaps → 6956 bytes per batch
- Total transactions in one **sequenceBatches**
 - Eth transfers $\Rightarrow (128k / 39630) * 350 = 1127$ txs
 - ERC20 transfers $\Rightarrow (128k / 24043) * 140 = 745$ txs
 - UniswapV2 swaps $\Rightarrow (128k / 6956) * 19 = 349$ txs
- Conversions: 1 ETH = 1900\$, gasPrice = 50 GWei

Assumptions \Rightarrow

Cost computation

sequenceBatches SC CALL ⇒ 119 002 GAS (constant cost)

Cost vs. Bytes



Cost computation

verifyBatches SC CALL

(get last 100 verifyBatches events)



- 351 920 GAS (constant cost)
- 273 batches verified in one single call

Single batch proof



- 1 batch proof computation cost¹: 0.0259 \$

¹<https://twitter.com/eduardiez/status/1667221787178876963?s=46&t=2Fzkb9sZOyy4irzxy70HVg>

Do the numbers

Ether transfer

Data-availability	→	<ul style="list-style-type: none">■ $157 \text{ bytes} \cdot 16 \text{ gas} = 2512 \text{ gas}$■ $2512 \text{ gas} \cdot 50 \text{ GWei} = 0.0001256 \text{ ether}$■ $0.0001256 \text{ ether} \cdot 1900 \\$ / \text{Eth} = \mathbf{0.2386 \\$}$
-------------------	---	---

Sequencing	→	<ul style="list-style-type: none">■ $119\,002 \text{ gas} / 1127 \text{ txs} = 106 \text{ gas}$■ $106 \text{ gas} \cdot 50 \text{ GWei} = 0.0000053 \text{ ether}$■ $0.0001256 \text{ ether} \cdot 1900 \\$ / \text{Eth} = \mathbf{0.01007 \\$}$
------------	---	---

Prover	→	<ul style="list-style-type: none">■ $0.0259 \\$ / 350 \text{ txs} = \mathbf{0.000074 \\$}$
--------	---	---

Aggregation	→	<ul style="list-style-type: none">■ $351\,920 \text{ gas} / (273 \text{ batches} \cdot 350 \text{ txs}) = 4 \text{ gas}$■ $4 \text{ gas} \cdot 50 \text{ GWei} = 0.0000002 \text{ ether}$■ $0.0000002 \text{ ether} \cdot 1900 \\$ / \text{Eth} = \mathbf{0.00038 \\$}$
-------------	---	--

INTRO: L2 transaction cost			
ERC20 transfer			
Data-availability	→	■	$201 \text{ bytes} \cdot 16 \text{ gas} = 3216 \text{ gas} \Rightarrow \mathbf{0.305 \$}$
Sequencing	→	■	$119\,002 \text{ gas} / 1127 \text{ txs} = 160 \text{ gas} \Rightarrow \mathbf{0.0152 \$}$
Prover	→	■	$0.0259 \$ / 140 \text{ txs} = \mathbf{0.000185 \$}$
Aggregation	→	■	$351\,920 \text{ gas} / (273 \text{ batches} \cdot 140 \text{ txs}) = 10 \text{ gas} \Rightarrow \mathbf{0.00095 \$}$
UniswapV2 swap			
Data-availability	→	■	$671 \text{ bytes} \cdot 16 \text{ gas} = 10736 \text{ gas} \Rightarrow \mathbf{1.019 \$}$
Sequencing	→	■	$119\,002 \text{ gas} / 349 \text{ txs} = 341 \text{ gas} \Rightarrow \mathbf{0.0323 \$}$
Prover	→	■	$0.0259 \$ / 19 \text{ txs} = \mathbf{0.001363 \$}$
Aggregation	→	■	$351\,920 \text{ gas} / (273 \text{ batches} \cdot 19 \text{ txs}) = 68 \text{ gas} \Rightarrow \mathbf{0.0064 \$}$

EIP 4844



- Directly reduce gas costs from 16 gas per byte to 3 gas per byte

- Reduce amount of bytes posted on-chain

Data compression



- Approach:
 - **Specific** transaction fields compression



2

zkBlob

New data structure



- Allow to use EIP4844 when added to Ethereum
- Add single verification polynomial commitment (EIP4844) or hash data (current approach)

Allows to



- Aggregate signatures for each batch
 - Add SNARK proof: all transactions signatures aggregated
 - SNARK proof verified in the zkBlob
- Decoupling decompression stage from the processing stage (Batch)
 - zkBlob to build transaction from compressed data
 - Encode parameters into a custom format to be ready for batches

Current approach

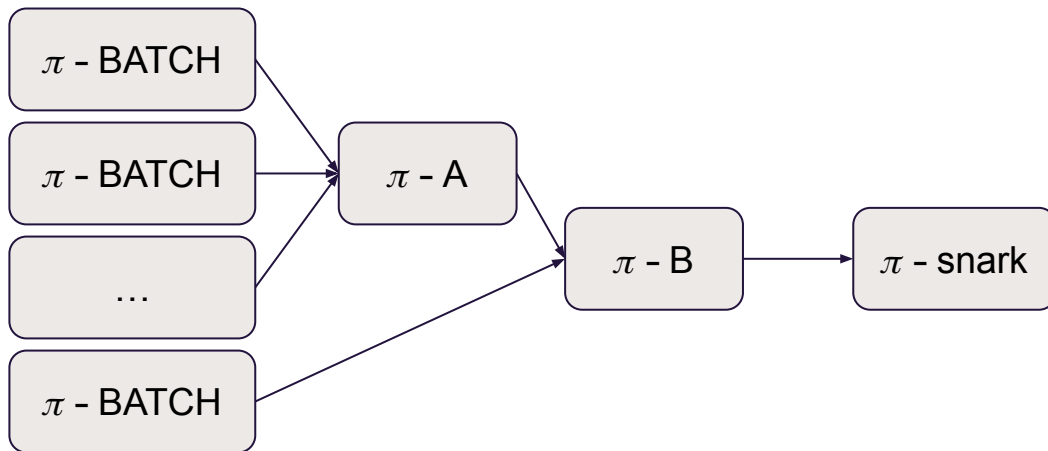
```
function sequenceBatches(BatchData[] calldata batches, address I2Coinbase)
```

zkevm-rom

BATCH

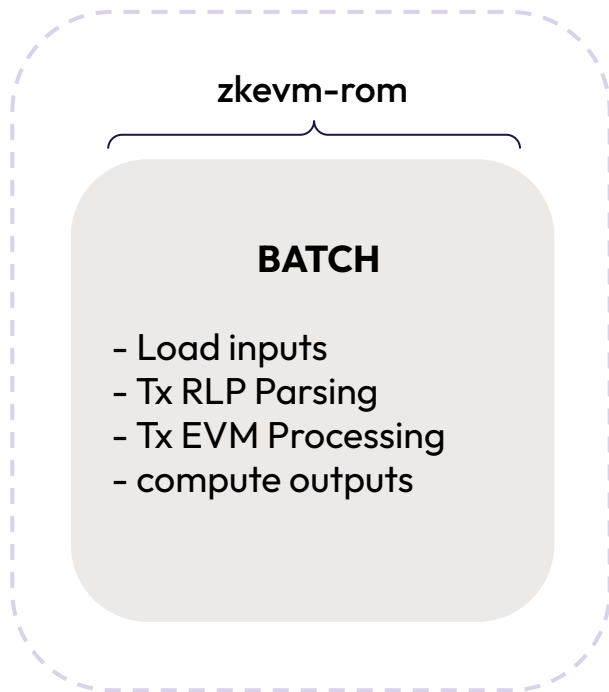
- Load inputs
- Tx RLP Parsing
- Tx EVM Processing
- compute outputs

Proof diagram



Current approach

```
function sequenceBatches(BatchData[] calldata batches, address I2Coinbase)
```



Add zkBlob

```
function sequenceZkBlob(bytes32 hashBlob)
```

zkevm-blob

- load inputs
- decompress tx data
- compute tx signed message
 - encode RLP tx
 - keccak(RLP[txFields])
 - ecrecover: retrieve **from**
- Build batches data
- Verify blob data (eip4844)
- Compute outputs

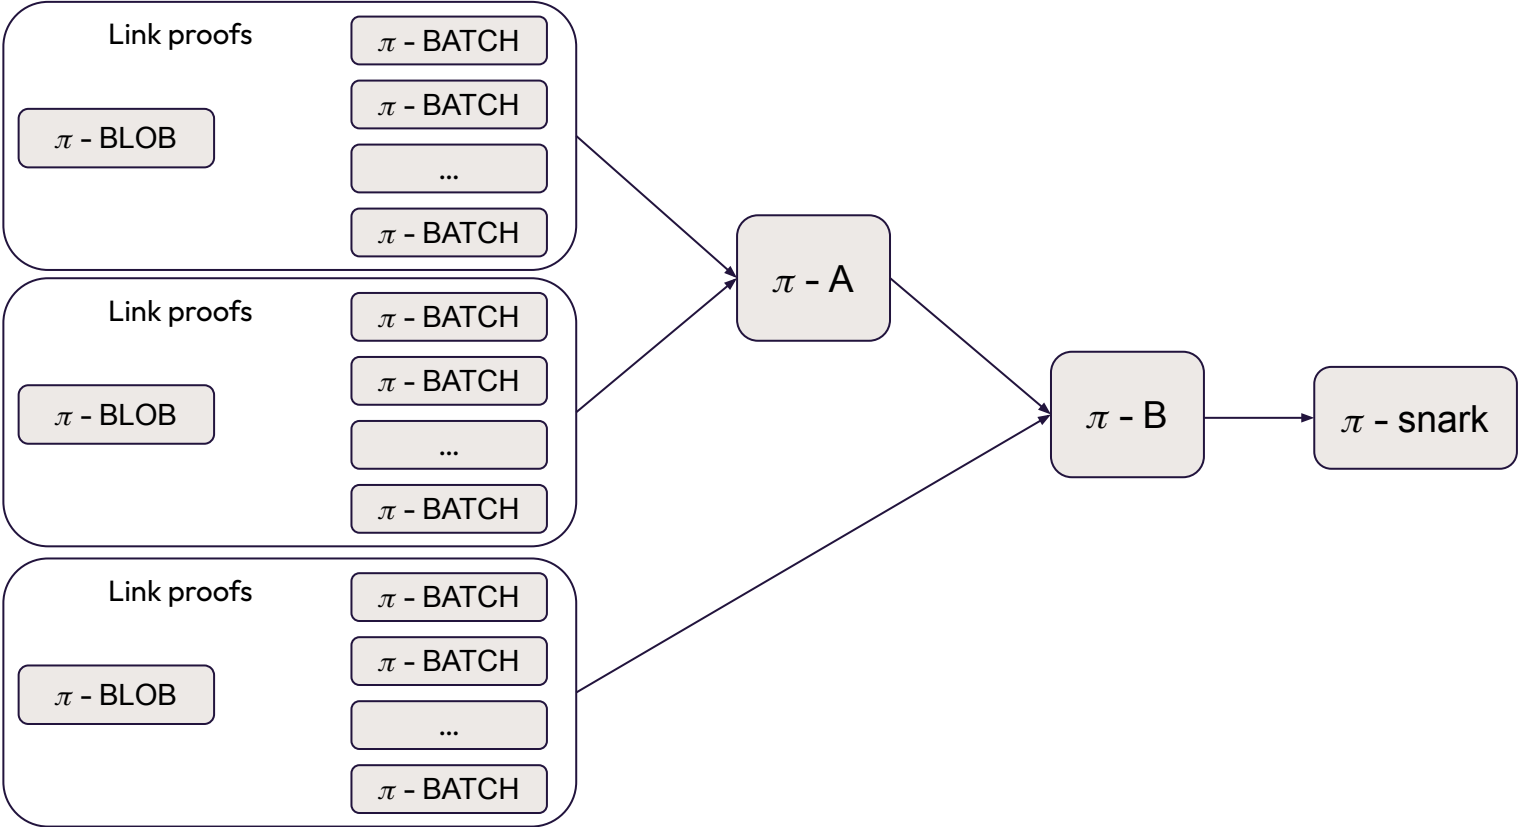
zkevm-rom

BATCH

- load input from zkBlob
- Tx EVM Processing
- compute outputs

Verify proof aggregation signatures

Proof diagram





3

Data compression

Remove signatures



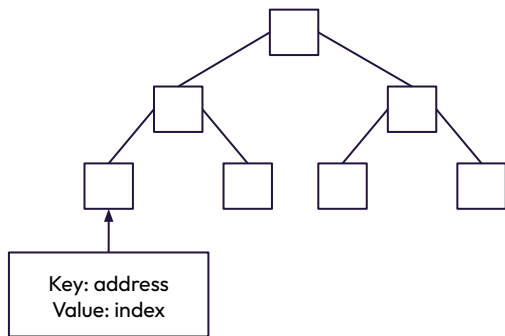
- Save 65 bytes (r, s, v) per transaction
- Cost data-availability SNARK proof shared at sequencing level by all transactions

Custom compression



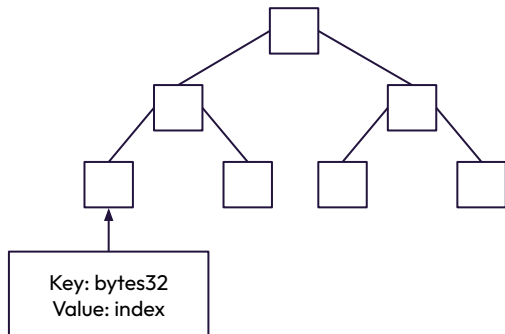
- Set of custom functions: $F(input_bytes) = output_bytes$
- Reconstruct transaction bytes from compressed bytes
- Two new SMT structures
 - Address tree (address alias)
 - Data tree (32 bytes alias)

Address tree



- Incremental index
- Forced by the circuit
- Alias addresses

Data tree



- Incremental index
- Limited size
- Forced by the circuit
- Alias 32 bytes
- Useful for **data** field (siblings merke trees)

Custom function set

- Small set of functions
- Specific behaviour if executed in *data* field
- 3 bits → function definition

First byte		Name: Data < 32 bytes
3 bits (header)	5 bits (header payload)	Description
000	XXXXX	Read next XXXX bytes (up to 32 bytes)

First byte		Name: Large data
3 bits (header)	5 bits (header payload)	Description
001	XXXXX	Read next XXXX bytes which is the length of the bytes to read

First byte		Name: Small value
3 bits (header)	5 bits (header payload)	Description
010	XXXXX	Value is coded in the body

DATA COMPRESSION: Specification

First byte		Name: 32 bytes compressed
3 bits (header)	5 bits (header payload)	Description
011	XXXXX	Read next XXXX bytes which will be the index in the Data tree

First byte		Name: address compressed
3 bits (header)	5 bits (header payload)	Description
100	XXXXX	Read next XXXX bytes which will be the index in the Address tree

First byte		Name: value compressed
3 bits (header)	5 bits (header payload)	Description
101	XXXXX	Read next XXXX bytes which will be the mantissa and 1 byte for the exponent ($V = m \cdot 10^e$)

DATA COMPRESSION: Specification

First byte		Name:uncompressed address
3 bits (header)	5 bits (header payload)	Description
110	00000	Read 20 bytes and store it in the address tree

First byte		Name: uncompressed 32 bytes
3 bits (header)	5 bits (header payload)	Description
110	00001	Read 32 bytes and add it in the data tree

First byte		Name: data < 32 bytes pad right
3 bits (header)	5 bits (header payload)	Description
111	XXXXX	Read next XXXX bytes and pad right zeros until 32 bytes

Extra Rule

- Pad left until 32 bytes if decompression is in **data** transaction field
 - small value, address/data un/compressed, value compressed

DATA COMPRESSION: Example

Transaction fields: ERC20 Transfer

nonce: 0

gasPrice: 10000000000

gasLimit: 100000

to: 0x1275fbb540c8efc58b812ba83b0d0b8b9917ae98

value: 0

data:

[illegible]

chainId: 1000

Non-compressed: **rlp(nonce, gasprice, gaslimit, to, value, data, chainId, 0, 0)**

[illegible]

Compressed: **txType # txBody** \Rightarrow **nonce, gasprice, gaslimit, to, value, data, chainId**

0x4140a10109a1010581044004a9059cbb810c01640203e8

Data availability gain

Data compression:

	Length (bytes)		
Tx Type	Message to hash + signature	New full encoding	Gain
ETH TRANSFER	111	15	7.40x
ERC20: TRANSFER	174	23	7.57x
UNISWAP: SWAP_EXACT_TOKENS_FOR_TOKENS	368	37	9.95x

EIP4844: 5x gain

Summary (applying reduction)

	\$ (%)			
	data-availability	sequencing	prover	aggregation
Ether transfer	0.0064 \$ (37.98 %)	0.010 \$ (59.35 %)	0.00007 \$ (0.42 %)	0.00038 \$ (2.26 %)
ERC20 transfer	0.0080 \$ (33.29 %)	0.015 \$ (62.03 %)	0.00018 \$ (0.74 %)	0.00095 \$ (3.93 %)
UniswapV2 swap	0.0204 \$ (33.91 %)	0.032 \$ (53.19 %)	0.00136 \$ (2.26 %)	0.00640 \$ (10.64 %)

	Eip4844 & data compression	
	before	after
Ether transfer	0.248 \$	0.016 \$
ERC20 transfer	0.321 \$	0.024 \$
UniswapV2 swap	1.058 \$	0.060 \$

Thanks !!

Carlos Matallana
Protocol & Integration Team Lead
carlos@polygon.technology

 @KrlsMata

